

# Deep Learning in *Computer Vision*

# What is Computer Vision

- The goal in computer vision is to extract useful information from images
- This includes, but is not limited to:
  - Reconstructing properties such as shape, illumination and color distributions
  - Detecting objects - e.g. faces in images
  - Estimating motion in image sequences
  - Detecting and matching points of special interest between images - e.g. for creating panorama images
  - ...

# Image classification



**Dog**

**Piano**

**Cake**

**Computer**

**Unknown**

# Image classification



Dog

Piano


Cake

Computer

Unknown



# Object Detection



# Segmentation





# Generative Models



# *Deep Learning* in Computer Vision



# Learning

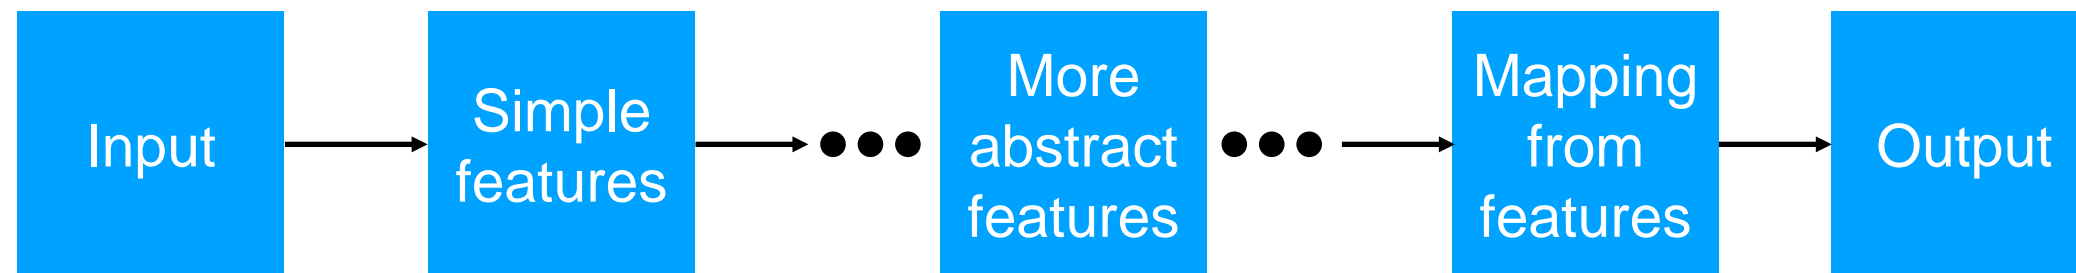
- Supervised learning
  - Learn mapping from input to output given training examples with known output
- Unsupervised learning
  - Learn “something” about the distribution of input examples without having known output
- Reinforcement learning
  - Learn actions based on rewards

# Deep Learning

- Deep learning uses deep neural networks
  - Neural networks can approximate any function, if they have enough parameters

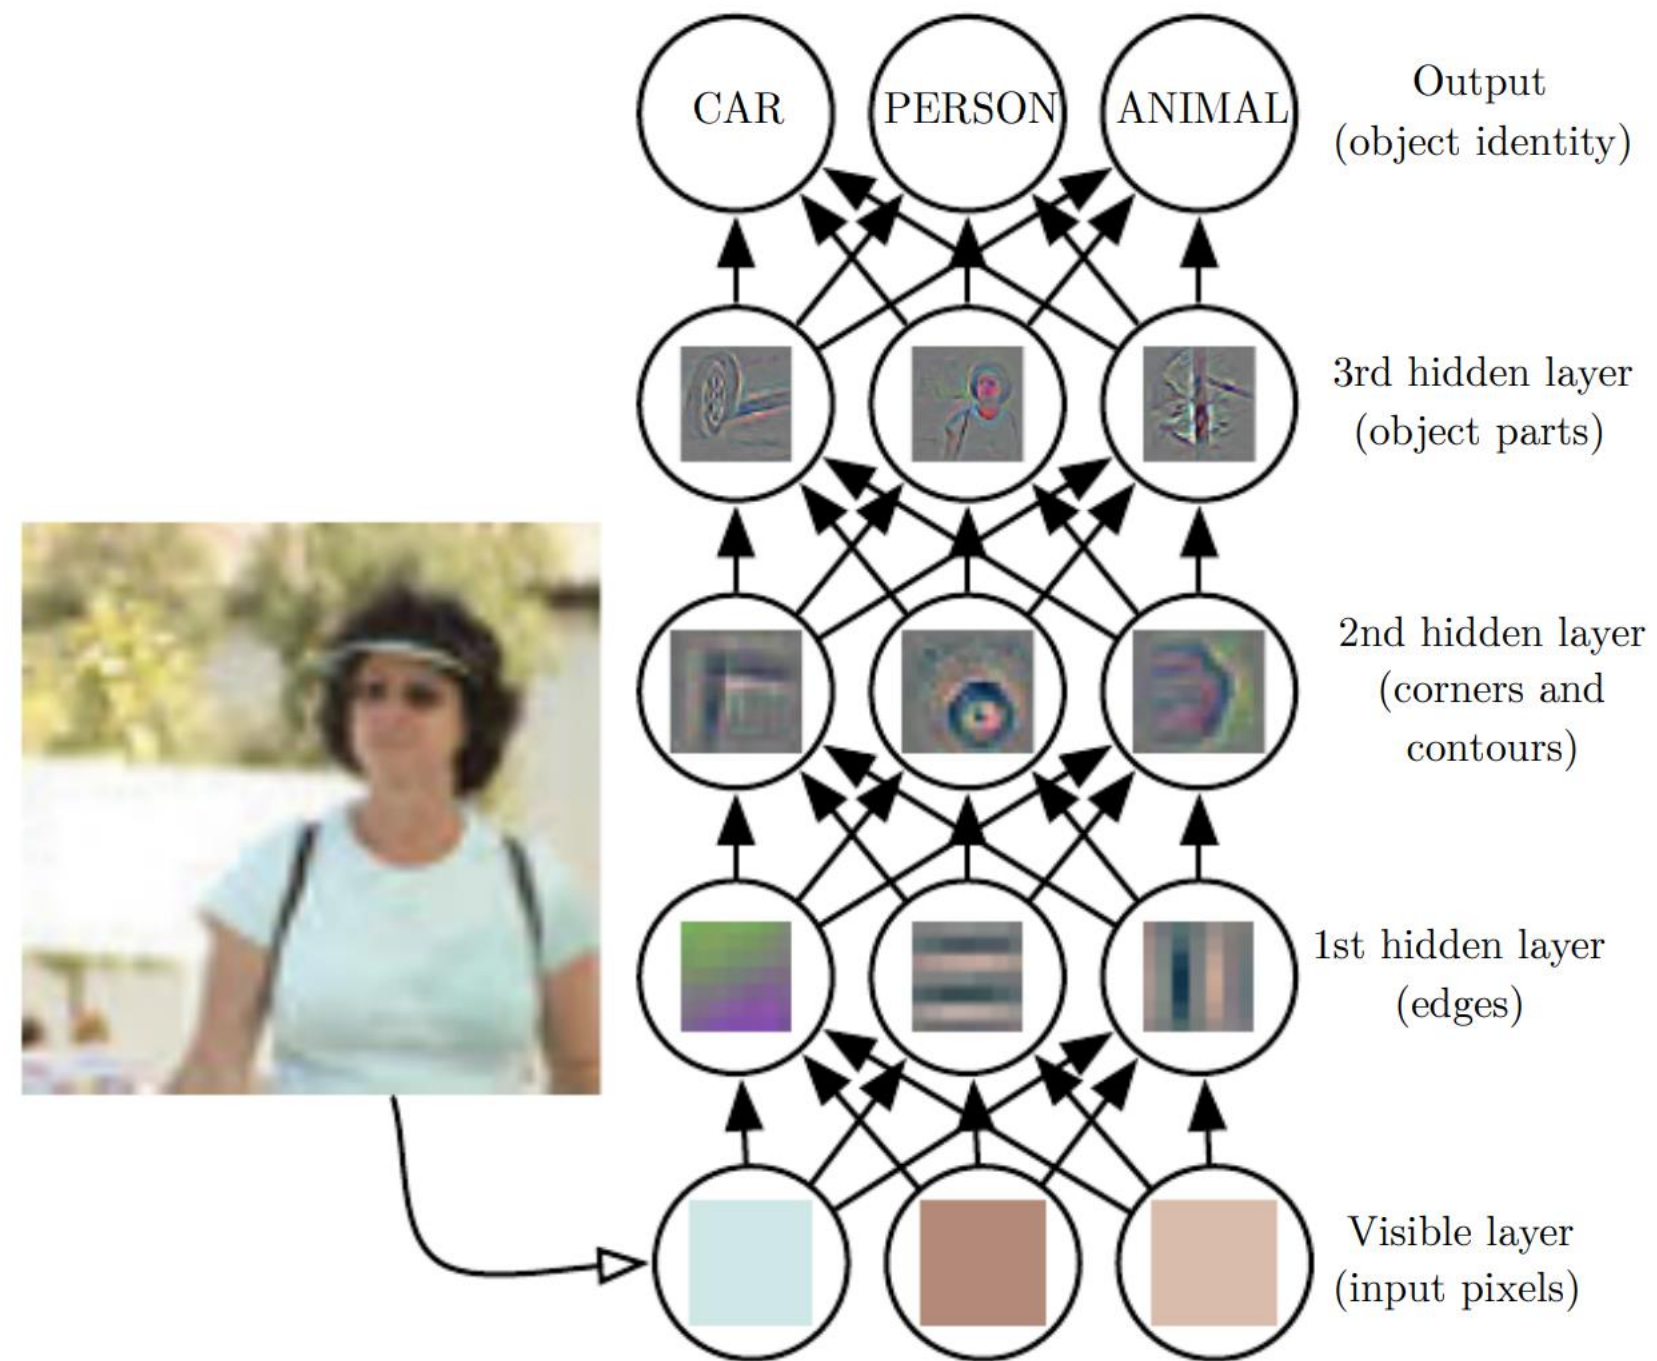
# Neural Networks

- In deep learning we seek to map a set of input values to output values
- Going directly from input to output is in most cases not possible
- Instead, we learn representations of the inputs from which it is easier to predict the output
- In deep learning we learn increasingly complex representations/features that are expressed in terms of simpler representations/features





# Neural Networks




# Neural networks

- The previous example can be written as:

$$\begin{array}{ccc} f(x) = f_4(f_3(f_2(f_1(x)))) \\ \uparrow & & \uparrow \\ \text{Output} & & \text{Input} \end{array}$$


- How do we represent the functions  $f_1, f_2, f_3, f_4$ ?

# Feed-Forward Neural networks







# Feed-Forward Neural networks



# Feed-Forward Neural networks




# Feed-Forward Neural networks






# Feed-Forward Neural networks



# Neurons - the building block

- Originally inspired by neurons in the brain




$$z = w_1x_1 + w_2x_2 + w_3x_3 + b = \mathbf{w}\mathbf{x} + b$$

$a = \sigma(z)$  is the output of the neuron

$\sigma(\cdot)$  is a non-linear activation function


# Feed-forward computation



$$\mathbf{a}^0 = \mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix}$$
$$\mathbf{a}^1 = \sigma(\mathbf{W}^1 \mathbf{a}^0 + \mathbf{b}^1)$$
$$= \sigma \left( \begin{bmatrix} w_{0,0}^1 & w_{0,1}^1 & w_{0,2}^1 \\ w_{1,0}^1 & w_{1,1}^1 & w_{1,2}^1 \\ w_{2,0}^1 & w_{2,1}^1 & w_{2,2}^1 \end{bmatrix} \begin{bmatrix} a_0^0 \\ a_1^0 \\ a_2^0 \end{bmatrix} + \begin{bmatrix} b_0^1 \\ b_1^1 \\ b_2^1 \end{bmatrix} \right)$$
$$\mathbf{a}^2 = \sigma(\mathbf{W}^2 \mathbf{a}^1 + \mathbf{b}^2)$$
$$\dots$$



# Feed-forward computation



$$\mathbf{a}^\ell = \sigma(\mathbf{W}^\ell \mathbf{a}^{\ell-1} + \mathbf{b}^\ell)$$


$$= \sigma \left( \sum_k w_{j,k}^\ell a_k^{\ell-1} + b_j^\ell \right)$$

$w_{j,k}^\ell$  is the weight from the  $k^{\text{th}}$  neuron in the  $(\ell-1)^{\text{th}}$  layer to the  $j^{\text{th}}$  neuron in the  $\ell^{\text{th}}$  layer

$b_j^\ell$  is the bias of the  $j^{\text{th}}$  neuron in the  $\ell^{\text{th}}$  layer

$a_j^\ell$  is the activation of the  $j^{\text{th}}$  neuron in the  $\ell^{\text{th}}$  layer

# Feed-forward computation



- We need a non-linear activation function after each layer
  - Otherwise, the entire network is a linear model (i.e., only one layer)

# How do we obtain the parameters?

- Right now, the model is just a lot of math, but how do we choose the weights and biases in the model, so it works well on our data?
- We need a function (loss) that measures how well our network is doing for a given set of parameters  $W$  and  $b$
- How can we use this to learn the parameters?

# How do we obtain the parameters?

- Random search?
  - Try many different random weights and biases and choose the best one
- Random local search?
  - Generate random weights and biases close to our current and choose the best one
- Gradient Descent
  - Use the gradient of our loss function, which gives the direction of maximum descent at any point  $x$

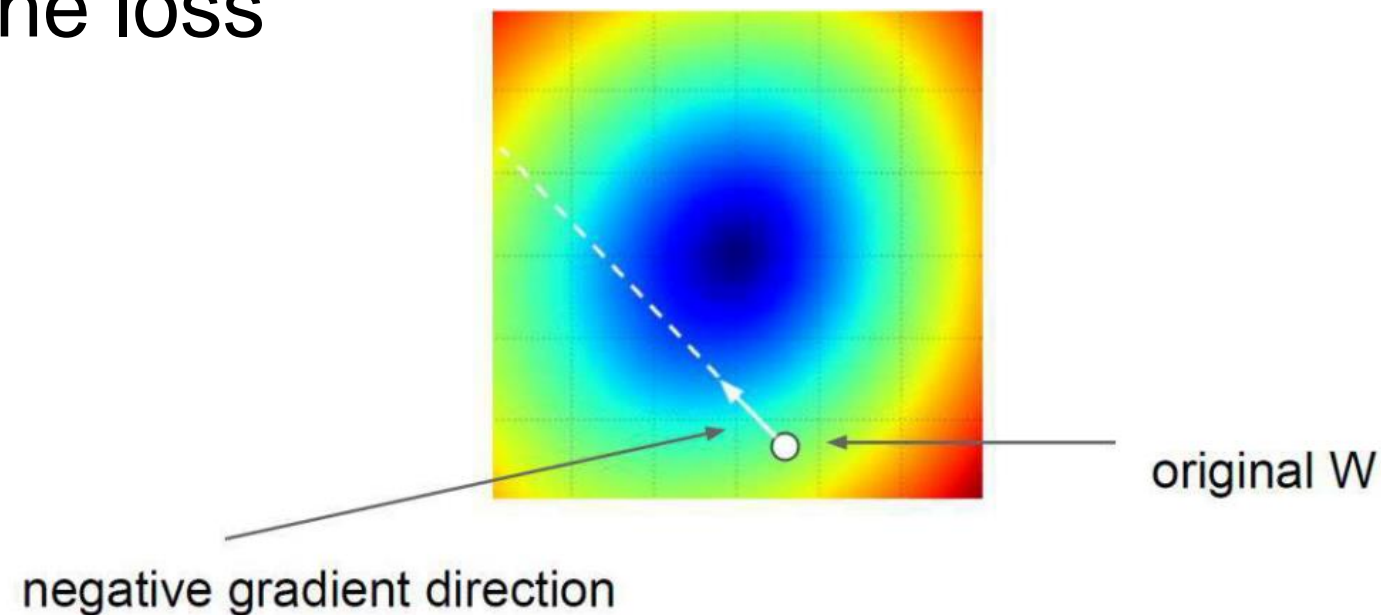
# How do we obtain the parameters?

- Random search?
  - Try many different random weights and biases and choose the best one
- Random local search?
  - Generate random weights and biases close to our current and choose the best one
- **Gradient Descent**
  - **Use the gradient of our loss function, which gives the direction of maximum descent at any point  $x$**



# Gradient descent

- **Gradient** descent
  - The gradient of a parameter, describes how the loss will change if we change this parameter
- Gradient **descent**
  - We change the value of the parameter slightly in the direction that should decrease the loss



# Gradient descent

- If  $\mathcal{L}$  is the loss function, we wish to minimize with respect to parameters  $p$ 
  - The gradient  $\nabla_p \mathcal{L}$  gives us the direction of maximum ascent of the loss function
  - $-\nabla_p \mathcal{L}$  is the direction of maximum descent
- Gradient descent algorithm:
  - Compute the gradient
  - Take a step in the negative gradient direction
$$p \rightarrow p' = p - \alpha \nabla_p \mathcal{L}$$
  - Here  $\alpha$  is called the learning rate and determines the step size
  - Repeat until convergence


# Backpropagation

- Which parameters do we want to find the partial derivatives of  $\mathcal{L}$  with respect to?
  - The weights  $\mathbf{W}$  and biases  $\mathbf{b}$
- How?
  - We propagate the error back through the network
    - Recursive approximation of the chain rule

# Short break

# Backpropagation example

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2x_2)}}$$




$f(x) = e^x$	→	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	→	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	→	$\frac{df}{dx} = a$		$f_c(x) = c + x$	→	$\frac{df}{dx} = 1$



# Backpropagation example


$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2x_2)}}$$



$f(x) = e^x$	$\rightarrow$	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	$\rightarrow$	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	$\rightarrow$	$\frac{df}{dx} = a$		$f_c(x) = c + x$	$\rightarrow$	$\frac{df}{dx} = 1$

# Backpropagation example

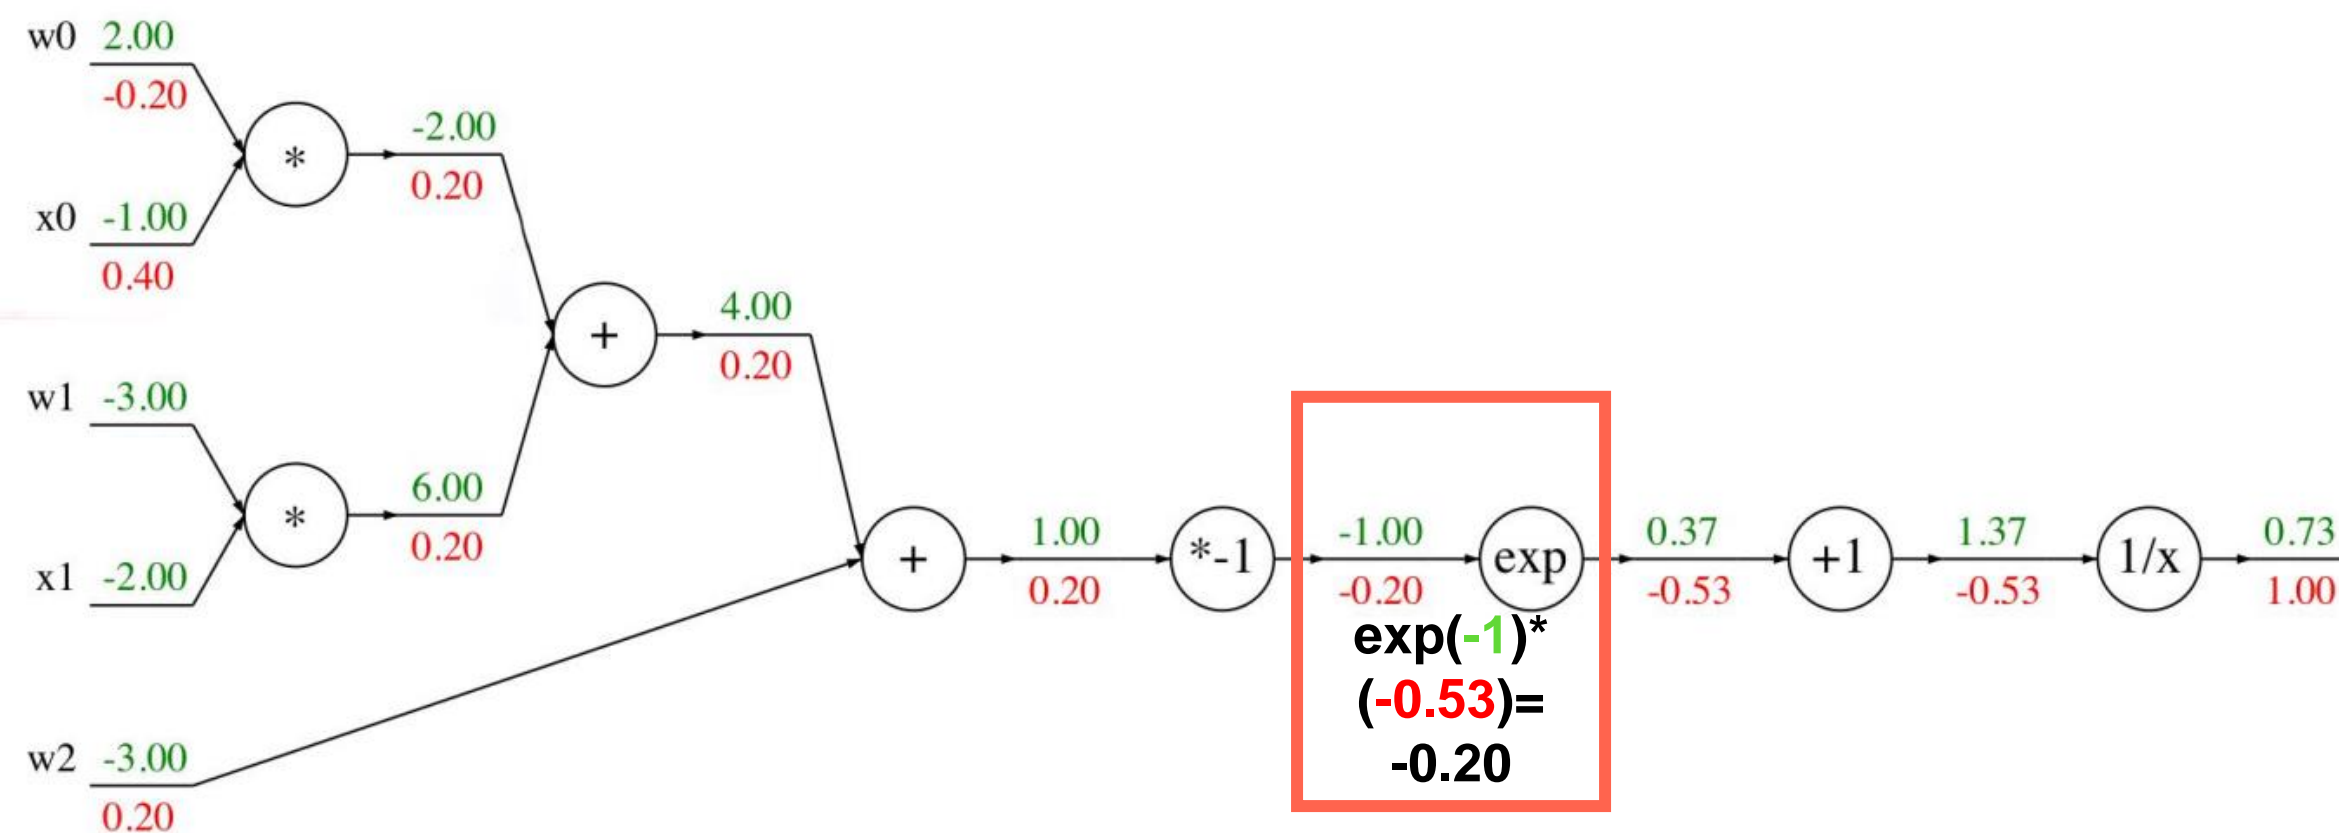
$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2x_2)}}$$



$f(x) = e^x$	$\rightarrow$	$\frac{df}{dx} = e^x$	$\left $	$f(x) = \frac{1}{x}$	$\rightarrow$	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	$\rightarrow$	$\frac{df}{dx} = a$	$\left $	$f_c(x) = c + x$	$\rightarrow$	$\frac{df}{dx} = 1$

# Backpropagation example


$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2x_2)}}$$



$f(x) = e^x$	$\rightarrow$	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	$\rightarrow$	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	$\rightarrow$	$\frac{df}{dx} = a$		$f_c(x) = c + x$	$\rightarrow$	$\frac{df}{dx} = 1$

# Backpropagation example

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2x_2)}}$$




$f(x) = e^x$	$\rightarrow$	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	$\rightarrow$	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	$\rightarrow$	$\frac{df}{dx} = a$		$f_c(x) = c + x$	$\rightarrow$	$\frac{df}{dx} = 1$




# Backpropagation example

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2x_2)}}$$



$$\begin{array}{lcl}
 f(x) = e^x & \rightarrow & \frac{df}{dx} = e^x \\
 f_a(x) = ax & \rightarrow & \frac{df}{dx} = a
 \end{array}
 \quad \Bigg| \quad
 \begin{array}{lcl}
 f(x) = \frac{1}{x} & \rightarrow & \frac{df}{dx} = -1/x^2 \\
 f_c(x) = c + x & \rightarrow & \frac{df}{dx} = 1
 \end{array}$$

# Backpropagation example




$f(x) = e^x$	$\rightarrow$	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	$\rightarrow$	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	$\rightarrow$	$\frac{df}{dx} = a$		$f_c(x) = c + x$	$\rightarrow$	$\frac{df}{dx} = 1$



# Backpropagation example

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2x_2)}}$$




$f(x) = e^x$	$\rightarrow$	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	$\rightarrow$	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	$\rightarrow$	$\frac{df}{dx} = a$		$f_c(x) = c + x$	$\rightarrow$	$\frac{df}{dx} = 1$

# Backpropagation example

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$$


What is the derivative here



$$\begin{array}{lcl}
 f(x) = e^x & \rightarrow & \frac{df}{dx} = e^x \\
 f_a(x) = ax & \rightarrow & \frac{df}{dx} = a
 \end{array}
 \Bigg|
 \begin{array}{lcl}
 f(x) = \frac{1}{x} & \rightarrow & \frac{df}{dx} = -1/x^2 \\
 f_c(x) = c + x & \rightarrow & \frac{df}{dx} = 1
 \end{array}$$

# Backpropagation example

$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2 x_2)}}$$



$$\begin{array}{lcl}
 f(x) = e^x & \rightarrow & \frac{df}{dx} = e^x \\
 f_a(x) = ax & \rightarrow & \frac{df}{dx} = a
 \end{array}
 \quad \Bigg| \quad
 \begin{array}{lcl}
 f(x) = \frac{1}{x} & \rightarrow & \frac{df}{dx} = -1/x^2 \\
 f_c(x) = c + x & \rightarrow & \frac{df}{dx} = 1
 \end{array}$$

# The equations of backpropagation

- Our goal is to find the partial derivatives of the loss function  $\mathcal{L}$  with respect to any weight  $w_{j,k}^\ell$  or bias  $b_j^\ell$

- The forward pass is defined by  $z_k^{\ell+1} = \sum_j w_{k,j}^{\ell+1} \sigma(z_j^\ell) + b_k^{\ell+1}$

- By the chain rule we have that 
$$\frac{\partial \mathcal{L}}{\partial w_{j,k}^\ell} = \frac{\partial \mathcal{L}}{\partial z_j^\ell} \frac{\partial z_j^\ell}{\partial w_{j,k}^\ell} = a_k^{l-1} \frac{\partial \mathcal{L}}{\partial z_j^\ell}$$
$$\frac{\partial \mathcal{L}}{\partial b_j^\ell} = \frac{\partial \mathcal{L}}{\partial z_j^\ell} \frac{\partial z_j^\ell}{\partial b_j^\ell} = \frac{\partial \mathcal{L}}{\partial z_j^\ell}$$

- We thus need a way to calculate  $\frac{\partial \mathcal{L}}{\partial z_j^\ell}$

# The equations of back propagation

- For the last layer in our network this is easily done ( $L$  is the last layer)

$$\frac{\partial \mathcal{L}}{\partial z_j^L} = \frac{\partial \mathcal{L}}{\partial a_j^L} \frac{\partial a_j^L}{\partial z_j^L} = \frac{\partial \mathcal{L}}{\partial a_j^L} \sigma'(z_j^L)$$

- The first term after the last equals sign depends on both the choice of loss function and choice of activation function

# The equations of back propagation

- Recall: 
$$z_k^{\ell+1} = \sum_j w_{k,j}^{\ell+1} \sigma(z_j^\ell) + b_k^{\ell+1}$$
- For the rest of the layers in the network we have  
( $\odot$  is elementwise product)

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial z_j^\ell} &= \sum_k \frac{\partial \mathcal{L}}{\partial z_k^{\ell+1}} \frac{\partial z_k^{\ell+1}}{\partial z_j^\ell} \\ &= \sum_k w_{k,j}^{\ell+1} \sigma'(z_j^\ell) \frac{\partial \mathcal{L}}{\partial z_k^{\ell+1}} \\ \frac{\partial \mathcal{L}}{\partial z^\ell} &= ((\mathbf{W}^{\ell+1})^T \frac{\partial \mathcal{L}}{\partial z^{\ell+1}}) \odot \sigma'(z^\ell) \end{aligned}$$



# Activation functions

- The two most commonly used activation functions are:

- Sigmoid  $\sigma(x) = \frac{1}{1 + e^{-x}}$

Closer to how the brain works



Vanishing gradient problem

- Rectified Linear Unit:  $\sigma(x) = \text{ReLU}(x) = \max(0, x)$

Works well in most cases

Not differentiable at zero

Not a problem in practice



# Derivatives of activation functions

- Sigmoid

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad \sigma'(x) = \sigma(x)(1 - \sigma(x))$$


- Rectified Linear Unit:

$$\sigma(x) = \text{ReLU}(x) = \max(0, x) \quad \sigma'(x) = \begin{cases} 1 & x > 0 \\ 0 & x \leq 0 \end{cases}$$

# Loss functions

- The choice of loss functions depends on the task
  - For regression: L2 (squared) or L1 (absolute)
  - For two-class classification: Binary cross-entropy
  - For multi-class classification: Cross-entropy
  - ...

# Two-class classification



- The output from the neuron in the last layer is mapped to the range  $[0,1]$  by using a sigmoid activation function on last neuron


# Binary cross-entropy

- If  $\hat{y} \in [0, 1]$  is the output from a neural network and  $y \in \{0, 1\}$  is the true value for an input  $x$  the binary cross-entropy is given by

$$\mathcal{L}(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

$$\frac{\partial \mathcal{L}}{\partial a^L} = \frac{\partial \mathcal{L}}{\partial \hat{y}} = -\frac{y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}}$$

# Multi-class classification



- The output from the neurons in the last layer is mapped to the range  $[0,1]$  such that they sum to 1 by using a softmax activation function in the last layer

# Softmax activation function

- The softmax activation function maps the outputs from all neurons in layer to the range  $[0,1]$  such that they sum to 1

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

$$\sum_i x_i = 1$$

- Softmax is shift invariant (if you add a constant to all  $x$ , the probabilities are the same)

# Cross-entropy

- The cross-entropy is given by

$$\mathcal{L}(\hat{y}, y) = - \sum_i y_i \log \hat{y}_i = - \log \hat{y}_I, \quad I = \arg_i(y_i = 1)$$

- i.e., the negative log likelihood
- Remember that to start the backpropagation we need to compute

$$\frac{\partial \mathcal{L}}{\partial z_j^L} = \frac{\partial \mathcal{L}}{\partial a_j^L} \frac{\partial a_j^L}{\partial z_j^L} = \frac{\partial \mathcal{L}}{\partial a_j^L} \sigma'(z_j^L)$$



# Updating the parameters

- The parameters can now be updated by taking a small step in the negative gradient direction

$$w_{j,k}^{\ell} \rightarrow \bar{w}_{j,k}^{\ell} = w_{j,k}^{\ell} - \alpha \frac{\partial L}{\partial w_{j,k}^{\ell}}$$

$$b_j^{\ell} \rightarrow \bar{b}_j^{\ell} = b_j^{\ell} - \alpha \frac{\partial L}{\partial b_j^{\ell}}$$

# Stochastic gradient descent

- So far, we only discussed how to train a network based on a single example
  - This is called *stochastic gradient descent*
- We can train over multiple training examples by simply averaging the loss and gradients over the examples
- If we use all our training examples, we call it *batch gradient descent*
- If we use random subsets of our examples, it is called *mini-batch gradient descent*
  - *Note: this is often called stochastic gradient descent (which is wrong, but widespread)*
- In practice we always use mini-batch gradient descent
  - The size of the minibatches is a hyperparameter
  - Typically chosen as large as RAM allows

# You have learned

- What neural networks are
- How to find the parameters for a neural network from data
  - Loss functions
  - Backpropagation

# Forming groups

If you do not have a group of 4 students, please go to the designated area and try to find one. Please try to align:

- Your level of ambition
- Your expected work hours

It is fine if you have different background/expertise. This can even be an advantage!