

Introduction to HPC

Dimitrios Papadopoulos
Associate Professor, DTU Compute

Deep Learning





Training and deploying a deep learning model

GPU is a must!!!

Options:

- Your CPU in your local machine (technically yes, but you will finish the assignment on July ☺)
- Google Colab (a free GPU, but not really...)
- The expensive GPU of your gaming desktop (technically yes, but you work in groups at campus)
- A GPU cluster (DTU has HPC cluster and we have allocated about 15 high-end GPUs (cost about 1M DKK) for this course)

Let's do a quick poll

(raise your hand if you agree)



(1) I am familiar with basic UNIX commands

*ls -l, pwd, cp, scp, rm -rf *, cd
. . . , mkdir, sudo, cat, rsync, df,
grep, vim, find, netstat*

(2) I know how to access a server via ssh

ssh user@awesomeserver.dtu.dk

(3) I have used colab or jupyter notebook

*jupyter notebook --no-browser --
port=9999*

(4) I have used pip, conda and/or virtualenv

```
conda create --name notmyfirstconda  
python=3.5
```

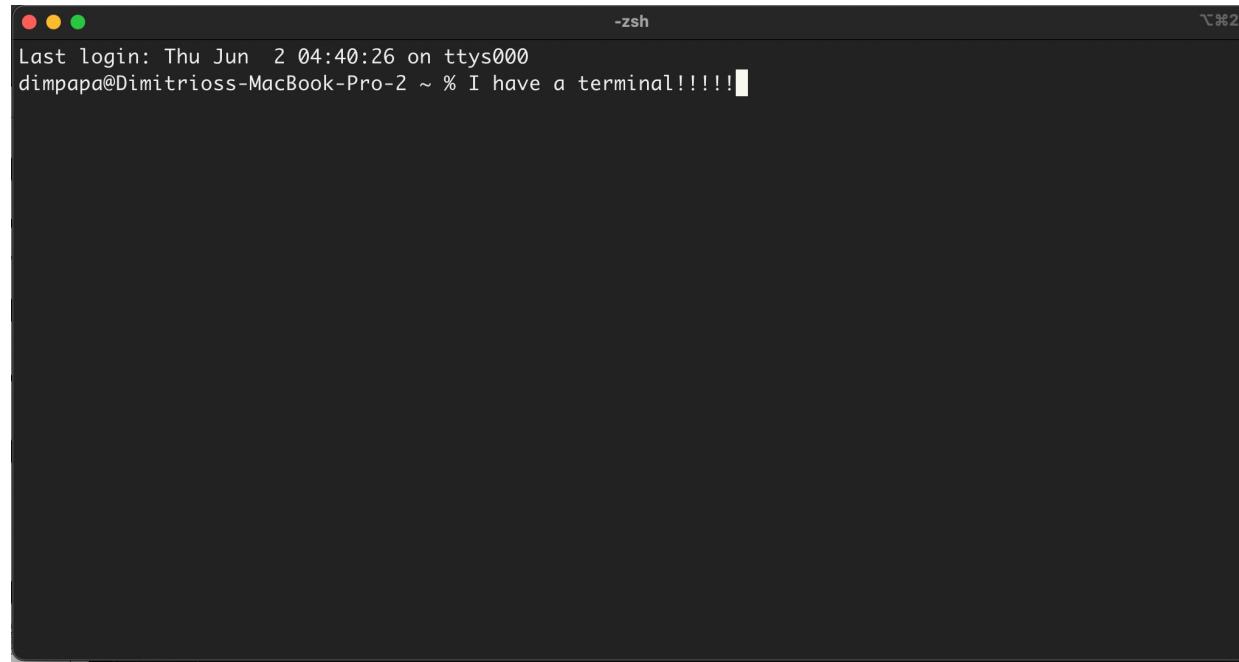
**(5) I am familiar with
running scripts on a cluster
(batch or interactive) using a
GPU**

Steps

- 1) Terminal
- 2) Learn Basic UNIX commands
- 3) Access HPC and open interactive node
- 4) Copy something to HPC filesystem
- 5) Monitoring commands and filesystem
- 6) Loading modules
- 7) Run a script (+GPU)
- 8) Conda or virtual environment
- 9) Open jupyter (graphics)
- 10) Open jupyter and forward it



(1) Terminal



Unix-based: default terminal or iTerm

Windows: default in new windows (?) or Putty, Kitty, windows terminal

(2) Basic UNIX commands

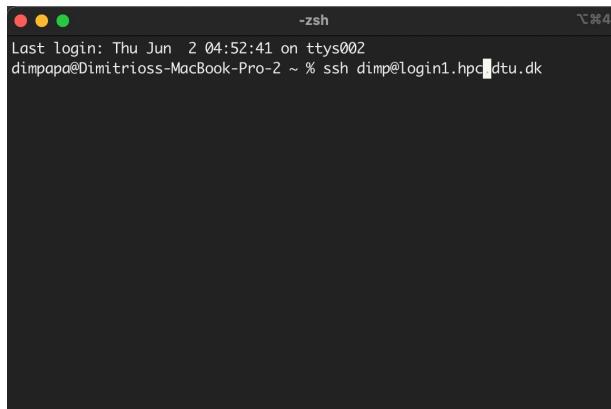
- Navigating to the filesystem
- Copy-Move-Create
- df/ du
- top/ htop
- man/ help
- nvidia-smi/ gpu-stat
- quota
- Googling for the rest if needed



<https://www.unixtutorial.org/basic-unix-commands>

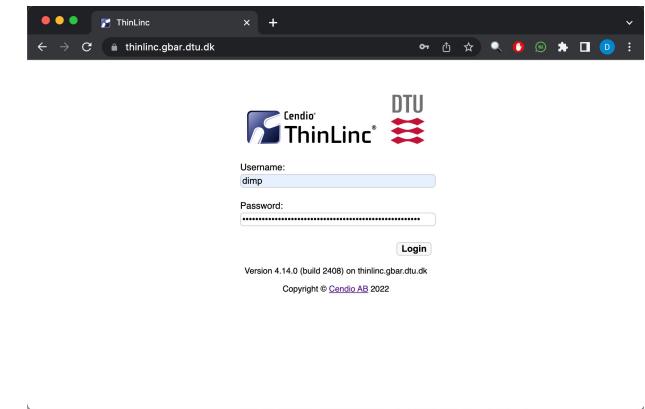
(3) Access HPC

- Everything about HPC: <https://www.hpc.dtu.dk/>
- Accessing the LSF 10 Cluster: https://www.hpc.dtu.dk/?page_id=2501



```
-zsh
Last login: Thu Jun  2 04:52:41 on ttys002
dimpapa@Dimitrioss-MacBook-Pro-2 ~ % ssh dimp@login1.hpc.dtu.dk
```

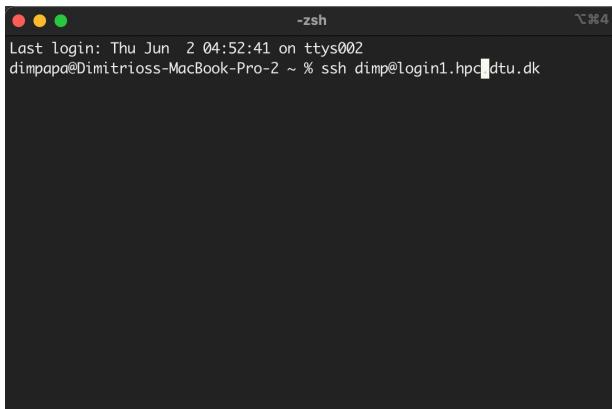
command line



thinlinc graphics

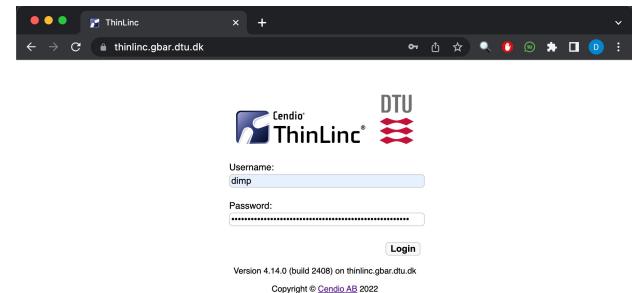
(3) Access HPC

- Everything about HPC: <https://www.hpc.dtu.dk/>
- Accessing the LSF 10 Cluster: https://www.hpc.dtu.dk/?page_id=2501



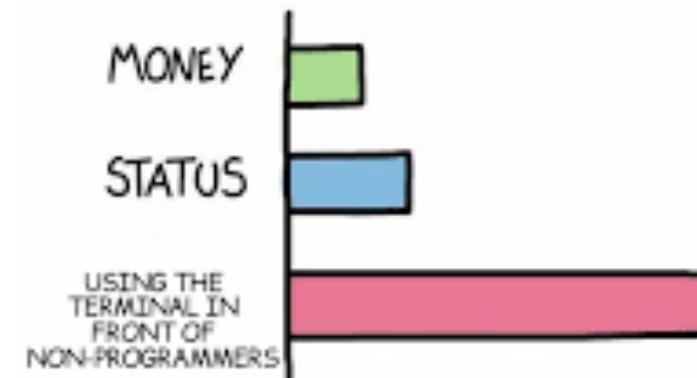
```
-zsh
Last login: Thu Jun  2 04:52:41 on ttys002
dimpapa@Dimitrioss-MacBook-Pro-2 ~ % ssh dimp@login1.hpc.dtu.dk
```

command line



thinlinc graphics

WHAT GIVES PEOPLE
FEELINGS OF POWER



(3) Access HPC

ssh-keys and VPN: https://www.hpc.dtu.dk/?page_id=4317

**When you
work at
the office**



**When you
work from
home**



When connecting from _outside_ of DTU (and not using DTU-VPN) one is required to authenticate with a ssh-key + ssh-key-keyphrase + DTU-password.

From inside of DTU (or connecting via DTU-VPN) one can also connect with password or ssh-key.

https://www.hpc.dtu.dk/?page_id=4317

(For thinlinc-users: This is not an error-message, this is an "information-message", please press "close" in the right-bottom-corner to continue.)

dimp@login1.hpc.dtu.dk: Permission denied (publickey).

(3) Access HPC

ssh-keys and VPN: https://www.hpc.dtu.dk/?page_id=4317



SSH and ThinLinc login using ssh-keys

To increase the security of the DTU infrastructure, the whole DTU network setup is changing, and this affects also the way users can access G-bar and the DCC HPC clusters **from outside the DTU network**.

Nothing changes, if you are accessing G-bar or the HPC cluster either on Campus, or via the **DTU VPN**. The instructions below are only relevant, if you need to access those resources from **outside** the DTU network, and the VPN is not a viable option, or it affects the performance critically (e.g. increased latency, incompatibility with a pre-existent pipeline, significant data transfer, etc).

If you only work on Campus or via the DTU VPN

Nothing changes for you. Just make sure that your **VPN** setup is up to date, and the **Multi Factor Authentication** is working.
You can safely skip the rest of the page.

If you work remotely, and always using the VPN is not a convenient option

SSH access will from now on require the use of ssh-keys + ssh-key passphrase + DTU password.

Instructions on how to setup your machine(s), and how to configure SSH for your G-bar/HPC account follow below.

Note: the guide below requires that you – during the setup process – are either connected to a DTU network on Campus, or via VPN (see DTU Inside for the **VPN setup**). Remember that the CISCO VPN requires the **Multi Factor Authentication**.

(3) Interactive nodes and Batch jobs

02514sh Is All You Need

Only 1 at a time →
if you can't open one →
kill the open session →
if not sure, *bkill XXXXX*

(4) Filesystem and copy

- **Monitoring:**

```
getquota_zhome.sh
```

- **Local machine to cluster:**

```
scp randomimage.png
```

```
dimp@login1.hpc.dtu.dk:/path/in/the/cluster/myfolder/bla/bla/
```

- **Shared space in cluster (read-only):**

```
/dtu/datasets1/02514/
```

(5) Monitoring

htop

```
x - o 
          0.0%] 6 [|
          0.0%] 11 [|
          0.0%] 16 [|
          0.0%] 0.0%] 7 [|
          0.0%] 12 [|
          0.0%] 17 [|
          0.0%] 0.0%] 8 [|
          0.0%] 13 [|
          0.0%] 18 [|
          0.0%] 0.0%] 9 [|
          0.0%] 14 [|
          0.0%] 19 [|
          0.0%] 0.0%] 10 [|
          0.0%] 15 [|
          0.0%] 20 [|
          0.0%]
Mem[|||||] 8.27G/125G Tasks: 161, 163 thr; 1 running
Swp[ 0K/20.0G Load average: 0.14 0.11 0.11
Uptime: 5 days, 05:34:22

PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command
3363 dimp 20 0 130M 2804 1592 R 1.3 0.0 0:00.10 htop
5253 s192270 20 0 1336M 55428 15412 S 0.7 0.0 12:02.53 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
1284 s213060 20 0 110M 1496 1232 S 0.0 0.0 0:22.93 bash -c while [ -d /proc/$PPID ]; do sleep 1;head -v -n 8 /proc/meminfo; head -v -
1707 root 20 0 579M 15376 4848 S 0.0 0.0 5:10.17 /usr/bin/python2 -s /usr/bin/fail2ban-server -xf start
2601 root 20 0 152M 10108 2668 S 0.0 0.0 0:12.47 /usr/bin/perl -wT /usr/sbin/munin-node --foreground
4993 s192270 20 0 3529M 2660M 18764 S 0.0 2.1 11:21.33 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
472 s192270 20 0 669M 180M 21616 S 0.0 0.1 3:33.73 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
1 root 20 0 187M 5096 2660 S 0.0 0.0 2:07.27 /usr/lib/systemd/systemd --switched-root --system --deserialize 22
319 root 20 0 180M 6380 4200 S 0.0 0.0 0:00.02 sshd: s192270 [priv]
356 s192270 20 0 183M 6268 1136 S 0.0 0.0 0:16.06 sshd: s192270@notty
371 s192270 20 0 110M 1732 1420 S 0.0 0.0 0:00.47 bash
473 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:00.00 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
474 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:02.04 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
475 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:02.00 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
476 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:01.85 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
477 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:01.81 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
478 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:00.00 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
479 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:19.13 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
480 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:19.19 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
481 s192270 20 0 669M 180M 21616 S 0.0 0.1 0:19.36 /zhome/23/5/144085/.vscode-server/bin/c3511e6c69bb39013c4a4b7b9566ec1ca73fc4d5/nod
F1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice F8Nice F9Kill F10Quit
```

nvidia-smi

```
n-62-20-1(dimp) $ nvidia-smi
Thu Jun  2 05:13:02 2022
+-----+
| NVIDIA-SMI 515.43.04    Driver Version: 515.43.04    CUDA Version: 11.7 |
|-----+-----+-----+
| GPU  Name      Persistence-MI Bus-Id     Disp.A | Volatile Uncorr. ECC | | | | |
| Fan  Temp  Perf  Pwr:Usage/Cap| Memory-Usage | GPU-Util  Compute M. |
| |          |          |          |           | GPU-Util  Memory M. |
|-----+-----+-----+
|  0  Tesla V100-PCIE...  On  | 00000000:37:00.0 Off |          0 | | | |
| N/A   31C   P0    25W / 250W | 0MiB / 16384MiB | 0%  Default |
|          |          |          |           |          | N/A |
+-----+-----+-----+
|  1  Tesla V100-PCIE...  On  | 00000000:AF:00.0 Off |          0 | | | |
| N/A   29C   P0    25W / 250W | 0MiB / 16384MiB | 0%  Default |
|          |          |          |           |          | N/A |
+-----+-----+-----+
| Processes:
| GPU  GI CI      PID  Type  Process name          GPU Memory |
| ID   ID          ID   ID          Usage          |
+-----+-----+-----+
| No running processes found
+-----+
```

getquota_zhome.sh

```
n-62-20-1(dimp) $ getquota_zhome.sh
You are using 817.51M of 30.00G quota
```

(Information updated every 30 minutes. Last Update: 2022-06-02 04:49:15)

Live demo Part 1!



(6) Modules

Modules: https://www.hpc.dtu.dk/?page_id=282

Home » HPC LSF clusters » LSF User Guides » Modules

Modules

- Why modules
- The module command
- Using modules
- Modules and batch jobs
- Loading modules automatically

Why modules

The choice of a modular approach to the software stack is natural, in a complex environment as HPC cluster. The reason behind modules is quite simple.

In a multi-user environment, especially when used by users from different disciplines, the software stack can become very large. However, each single user is only interested in a bunch of programs, not to all of them. Moreover, different users need sometimes different versions of the same software, and so these different versions should co-exist in the system. This may lead to potential system-wide version-conflicts, and to a confusion for the users, that may inadvertently use the wrong version.

A modular approach allows to have the full pool of software packages available to all of the users, and to simplify the management.

Using modules, each user has control over his own environment, for example the shell environment variable PATH, that holds the name of the directories from where you are allowed to run programs, without having to specify the full path. As a minor drawback, the user will not find all the programs he needs readily available at the login, but the program modules need to be loaded explicitly.

(7) Run a script

```
python i_can_run_hello_world_remotely.py
```

```
CUDA_VISIBLE_DEVICES=1 python i_can_even_run_it_on_a_gpu_that_i_select_wow.py
```

(8) Conda or Virtualenv

Virtual environments

Using virtual environments lets you have different stacks of software packages in different folders. Please [search the web](#) for additional information regarding virtual environments.

First, select the Python version you want, and load the corresponding module. Here we choose python 3.8.2. If you use versions older than 3.6, the instructions to create the virtual environment may differ.

```
module load python3/3.8.2
```

Then create a virtual environment. We choose the name `venv_1`

```
python3 -m venv venv_1
```

That's all. Now, activate the environment with the command

```
source venv_1/bin/activate
```

When the environment is active the shell notes this to the left of the cursor. It will look something like this:

```
(venv_1) ~  
gbarlogin1$
```

This indicates the virtual environment name in parenthesis.

Note that *both* the commands `python` and `python3` point to the python version of your original module. Now, when the environment is active, you can install all the packages you need with the normal `python3 -m pip install ...` syntax, for example (note `--user` should **NOT** be added)

```
python -m pip install scikit-learn
```

All those packages will be installed locally under the directory

```
venv_1/lib/python3.8/site-packages/
```

When you have finished to use the environment, just deactivate it with the command

```
deactivate
```

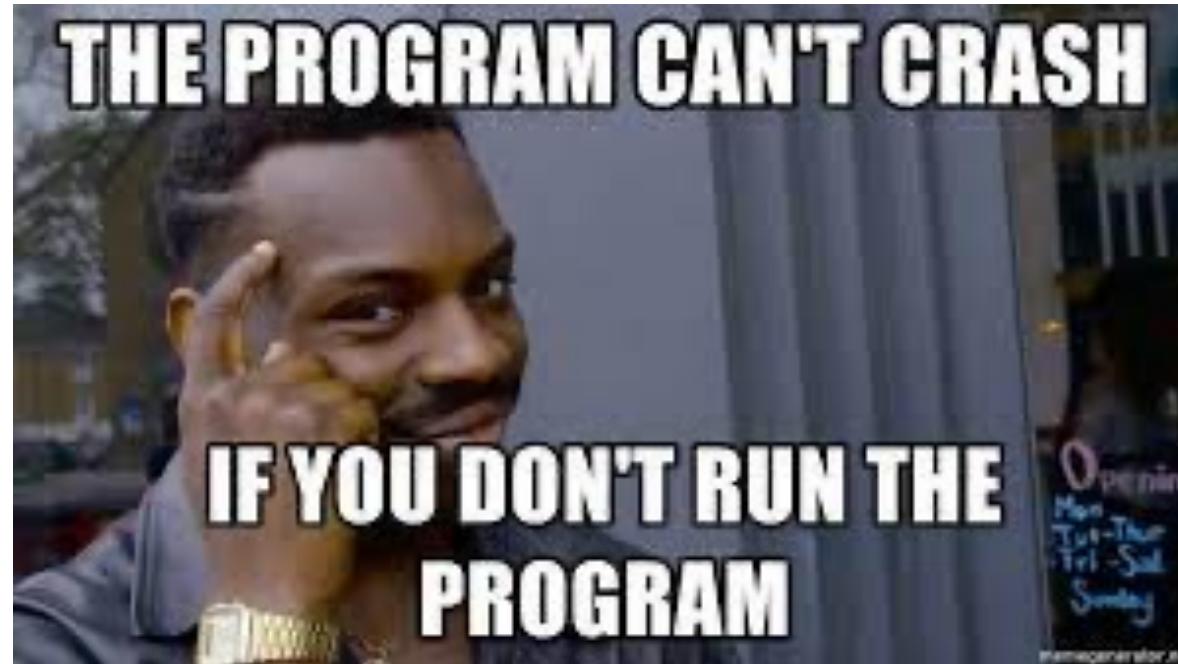
https://www.hpc.dtu.dk/?page_id=3678



(9) Jupyter

- A) jupyter notebook --no-browser --port=XXXXX --ip=\$HOSTNAME
- B) browser thinlinc: n00-00-0:XXXXX and token or password

It seems the easiest way. We don't recommend it at all. It will create you several problem



(10) Jupyter and forward port to my machine

Very handy command to ssh tunneling (port forwarding)!!!

Terminal 1 (local):

- (1) ssh to hpc: ssh **USER@login1.hpc.dtu.dk**
 - (2) open interactive node: **02514h**
 - (3) run jupyter: **jupyter notebook --no-browser --port=XXXXX --ip=\$HOSTNAME**

Terminal 2 (local):

- (1) port forwarding: **ssh USER@login1.hpc.dtu.dk -g -LXXXXXX:n-00-00-00:XXXXX -N**

Browser (local):

- (1) copy paste command from Terminal 1:

http://127.0.0.1:XXXXX/?token=xy

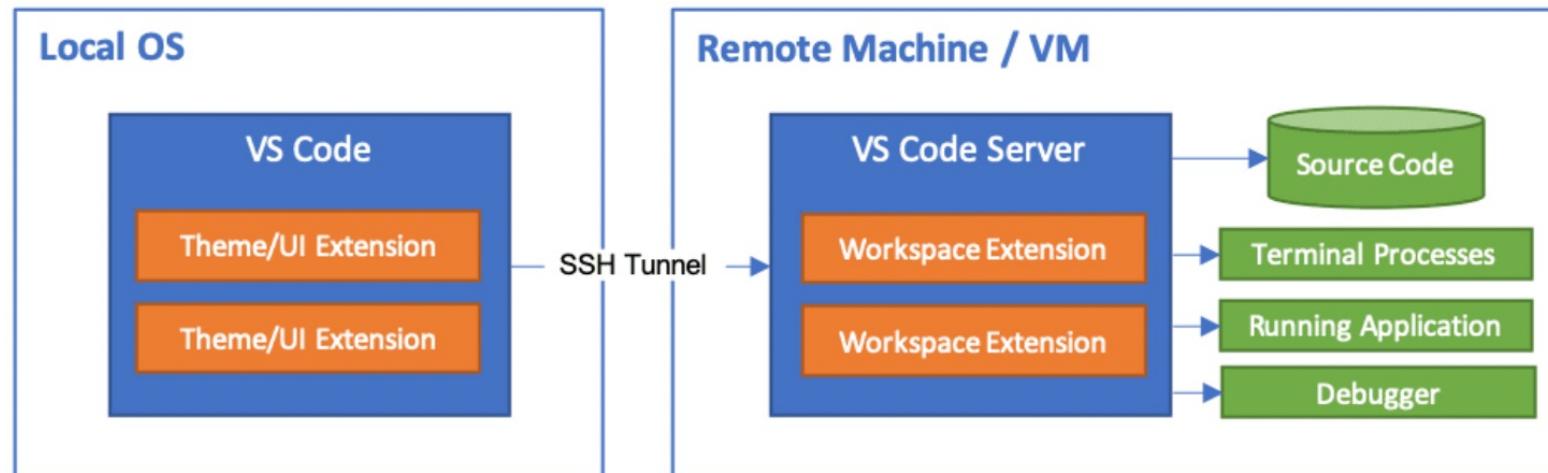
Live demo Part 2!



Bonus 1!!

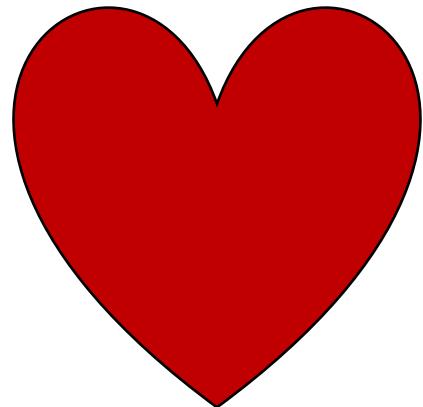
**Scripts / Classes/ Functions >>> Notebooks
BUT
Editors>>>vim**

Vscode and Vscode with remote ssh:
<https://code.visualstudio.com/docs/remote/ssh>

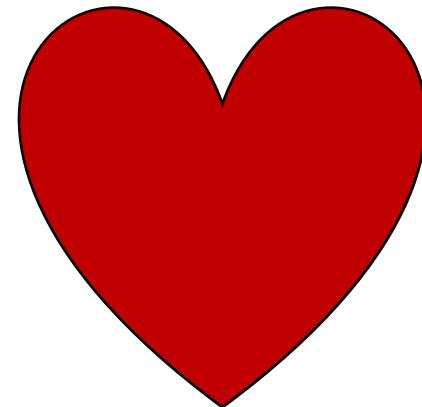


Bonus 2!!

“My terminal is awesome!”



tmux



Questions???

Exercise 1.1

11.00-12.00

Good Luck!