DLinCV Lecture 2.3

Image Segmentation: Poster feedback, small comment, and Q&A

12.06.2023

Aasa Feragen
afhar@dtu.dk

Poster feedback

A bit of general feedback from last week's poster session

- You can limit text by using short bullet points
- Great this time – keep for later: Put technical details on poster
- Giving each other feedback on posters – how did this go? Should it be more structured?

Dataset bias and learning spatial information

Why is the following excerpt from the Segment Anything paper relevant?

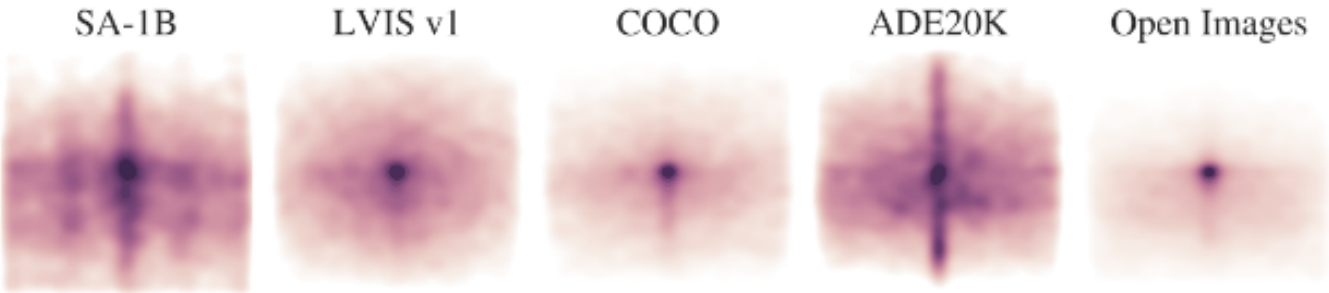

Figure 5: Image-size normalized mask center distributions.

Mask properties. In Fig. 5 we plot the spatial distribution of object centers in SA-1B compared to the largest existing segmentation datasets. Common photographer biases are present in all datasets. We observe that SA-1B has greater coverage of image corners compared to LVIS v1 [44] and ADE20K [117], the two most similarly distributed datasets, while COCO [66] and Open Images V5 [60] have a more prominent center bias. In Fig. 6 (legend) we compare these

Dataset bias and learning spatial information

Discuss: Can the original U-net learn spatial informtion? Can your implemented U-nets learn spatial information? Is this problematic?

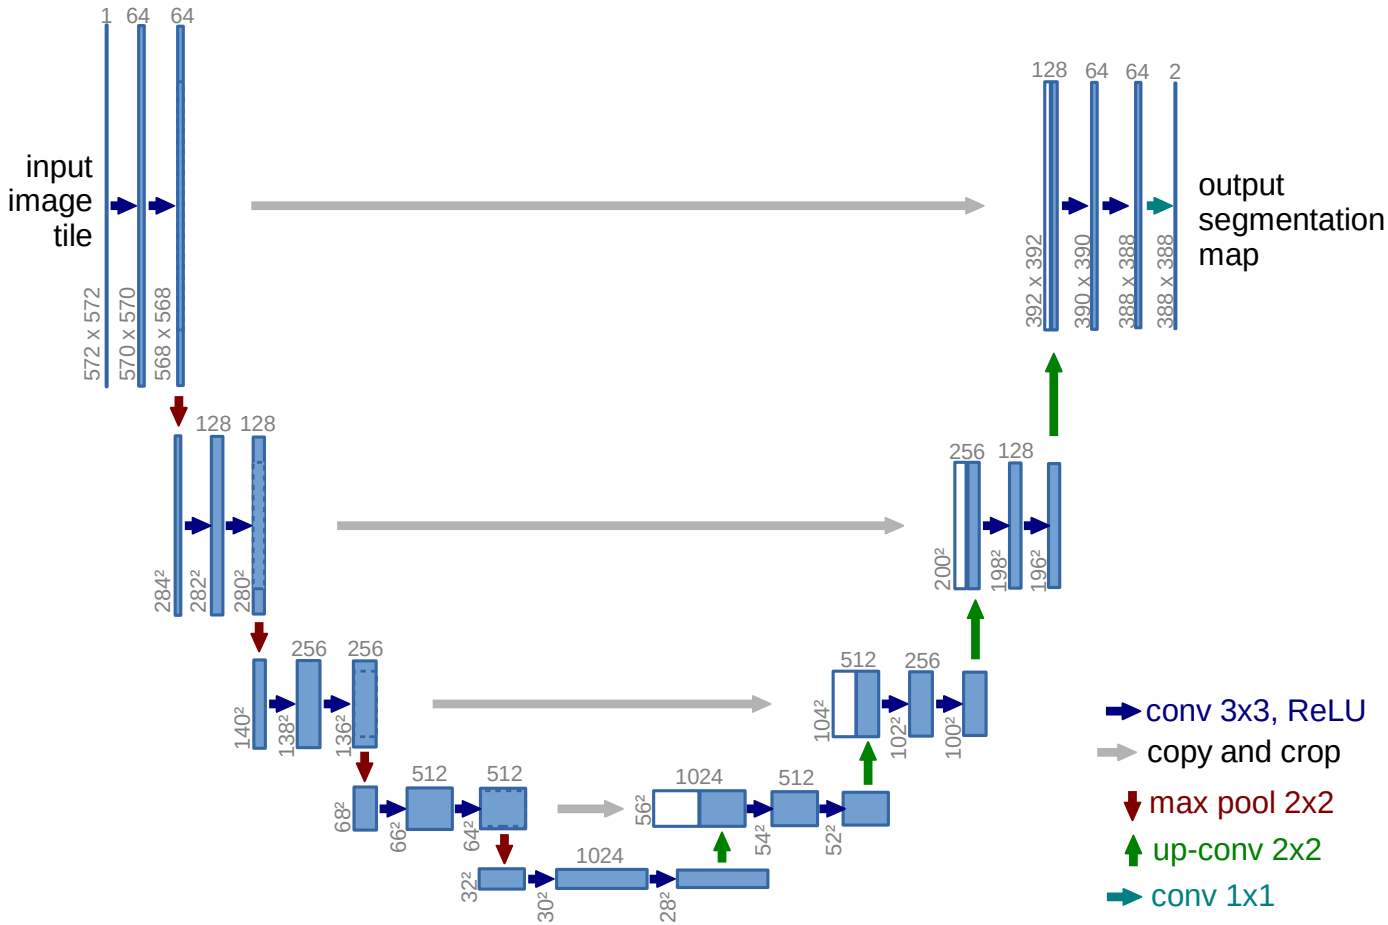

Dataset bias and learning spatial information

Discuss: Can the original U-net learn spatial informtion? Can your implemented U-nets learn spatial information? Is this problematic?

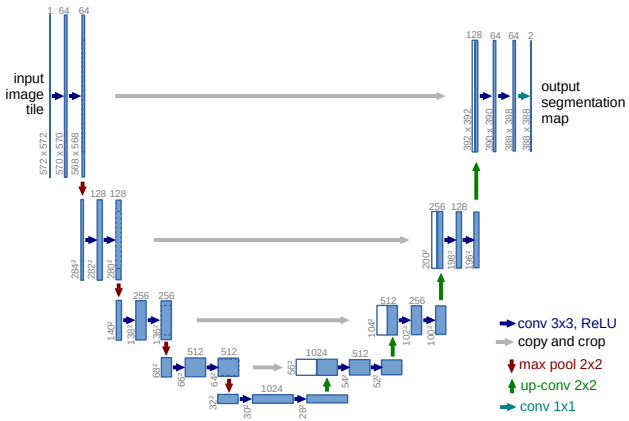

If you use zero padding, you implicitly add a black border around the image – the boundary has distinctively different features from the rest of the image – the network can learn how close to the boundary you are.

Dataset bias and learning spatial information

What about this one?

Q&A

How is it going?
Any questions/problems?